


Machine Learning Assignment 1

Name: Anh Quoc Do

NetID: aqd14

Part A


1. Plot of raw data



2. Converged value of θ

- First iteration:
 - $\theta = [0.0787538, 0.41494108]$
- Converged after 803 iterations (set error threshold $1e-5$)
 - $\theta = [0.49311086, 0.12820581]$


3. Plot of line that you fit




4. Prediction of your model for the height of two boys, age 3.5 and 7

- prediction = [0.94183121, 1.39055157]

5. Plot of 3D surface of J



6. Plot of contour of J



7. Answer to question about relationship between these plots and θ value found by your algorithm.


- These plots offer a visual view of how different values of θ affect to the value of cost function.
- There is exist an optimal vector θ that helps the cost function reaches minimum value. With gradient descent, in each step we update theta values that move toward optimal values

8. Include your code solution to Task A in your pdf submission

See appendix

Part B

1. Plot(s) comparing different learning rates



2. Converged value of θ for best learning rate

- $\theta = [353178.61702128, 114973.94457945, -3702.28587976]$ with $\alpha = 0.3$

3. Prediction of your model for the price of a house with 1650 square feet and 3 bedrooms

- Predicted price = \$303,614

4. Value of θ computed using normal equations

- $\theta = [109185.1639846, 123.62723839, -3185.43247293]$
- These values are different from those using gradient descent

5. Prediction of your normal equations θ on a house with 1650 square feet and 3 bedrooms.


- Predicted price = \$303,614

6. Include your code solution to Task b in your pdf submission

See Appendix


Part C

1. Plot of raw data



2. What values of θ did you get? How many iterations were required for convergence?
 - Converged at epoch 6
 - $\theta = [-15.76822471, 0.14087783, 0.15219571]$
3. What is the probability that a student with a score of 20 on Exam 1 and a score of 80 on Exam 2 will not be admitted?
 - $p = 0.684604$

4. Plot of final decision boundary found



Appendix

Part A & B

```
# Implmentation of linear regession algorithm for only one feature dataset  
# by using gradient descent to find optimum weights to the classification
```

```
'''
```

```
The data files ax.dat and ay.dat contain some example measurements of  
heights for various boys between the ages of two and eights. The y-values  
are the heights measured in meters, and the x-values are the ages of the boys  
corresponding to the heights.
```

```

There are 50 training examples in total
'''

# usage python linear_regression.py

import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
from numpy.linalg import inv

class LinearRegression(object):
    """ Linear regression model

    Parameters
    -----
    alpha : float
        Learning rate

    iter : integer
        Number of iteration

    check_conv: boolean
        Check if model is check_conv or not to stop iterations
    """

    def __init__(self, alpha=0.01, iter=2000, check_conv=True):
        self.alpha = alpha
        self.iter = iter
        self.check_conv = check_conv
        # self.normalized = normalized

    def fit(self, X, y):
        """ Fit traning dataset to the model.
        Update theta simultaneously by using gradient descent.
        For simplicity, ignore error handling.

        X : integer
            training input
        y : integer
            label output
        """

        self.theta = np.zeros((X.shape[1], 1)) # Column vector

```

```

err_threshold = 1e-5 # Stop when the cost function is smaller
than threshold
J = []
for i in range(self.iter):
    cost = cost_function(X, y, self.theta)
    J.append(cost)

    gra, _ = self.gradient(X, y)
    temp = self.theta - self.alpha * gra
    # if np.array_equal(temp, self.theta):
    if self.check_conv is True and np.sum(abs(self.theta -
temp)) < err_threshold:
        break
    # Simultaneously update theta
    self.theta = temp

    # print('Not convergered!')
    return self.theta, J

def hypothesis(self, X):
    # return np.dot(X, self.theta[1:]) + self.theta[0]
    return np.dot(X, self.theta)

def gradient(self, X, y):
    """Calculate gradient descent

    Parameters
    -----
    X : array-like
        training data

    Y : vector
        label output

    Returns
    -----
    gradient : float
        gradient value of given data

    mse : float
        mean squared errors
    """
    err = y - self.hypothesis(X)
    mse = (1.0/X.shape[0]) * np.sum(np.power(err, 2))

```



```

        gradient = -(1.0/X.shape[0]) * X.T.dot(err)
        return gradient, mse

    def predict(self, X):
        return self.hypothesis(X)

def cost_function(X, y, theta):
    """
    Calculate cost function
    """
    m = X.shape[0] # Training size
    err = y - np.dot(X, theta)
    J = 1.0/(2*m) * np.sum(np.power(err, 2))

    return J

# ----- PART A ----- #
def part_a():

    # load data
    X = np.loadtxt('data/ax.dat') # input data
    y = np.loadtxt('data/ay.dat') # output data

    # plot dataset
    plt.scatter(X, y, facecolors='blue')
    plt.xlabel('Age in years')
    plt.ylabel('Height in meters')
    # plt.show()

    # Data preprocessing
    m = X.shape[0] # Training size
    X = np.stack((np.ones(m), X), axis=-1)
    y = y.reshape(m, 1)

    model = LinearRegression(alpha=0.07)
    theta, J = model.fit(X, y)

    # Plot straight line
    plt.plot(X[:,1], np.dot(X, theta))
    plt.legend(['Linear Regression', 'Training Data'])
    plt.show()
    # plt.savefig('plot1.png')

    # Prediction
    test_data = np.array([[1, 3.5], [1, 7]])

```

```

prediction = model.predict(test_data)
print("Predicted height of kids age 3.5 and 7: ", prediction)

plot_surface_contour(X, y)

def plot_surface_contour(X, y):
    # Display Surface Plot of J
    t0 = np.linspace(-3, 3, 100).reshape(100, 1)
    t1 = np.linspace(-1, 1, 100).reshape(100, 1)

    T0, T1 = np.meshgrid(t0, t1)

    J_vals = np.zeros((len(t0), len(t1)))

    for i in range(len(t0)):
        for j in range(len(t1)):
            t = np.hstack([t0[i], t1[j]])
            J_vals[i, j] = cost_function(X, y, t)

    #Because of the way meshgrids work with plotting surfaces
    #we need to transpose J to show it correctly
    J_vals = J_vals.T

    # print(J_vals)

    fig = plt.figure()
    ax = fig.gca(projection='3d')
    ax.plot_surface(T0, T1, J_vals)
    plt.show()
    plt.close()
    #Display Contour Plot of J
    plt.contour(T0, T1, J_vals, np.logspace(-2, 2, 15))
    plt.show()

# ----- END PART A ----- #

# ----- PART B ----- #
def part_b():
    # Load data from files
    X = np.loadtxt('data/bx.dat')
    y = np.loadtxt('data/by.dat')

    sigma = np.std(X, axis=0) # std

```

```

mu = np.mean(X, axis=0) # mean

# Data preprocessing
m = X.shape[0]
y = y.reshape(m, 1)

X_scaled = (X-mu) / sigma # normalize(X)
intercept = np.ones((m, 1))
X_scaled = np.column_stack((intercept, X_scaled))
# print(X_scaled)
# y_scaled = normalize(y)

iterations = 50

alphas = [0.01, 0.03, 0.1, 0.3]
colors = ['b', 'r', 'g', 'y']

test_data = np.array([[1, 1650], [1, 3]])

for i in range(len(alphas)):
    model = LinearRegression(alpha=alphas[i], iter=iterations,
check_conv=False)
    theta, J = model.fit(X_scaled, y)

    # print('Learning rate %f' % alphas[i])
    # print('Theta: %d -- %d' % (theta[0], theta[1]))
    # print('Predict for the price of a house with 1650 square
feet and 3 bedrooms: ', model.predict(test_data))

    # Now plot J
    plt.plot(range(iterations), J, color=colors[i],
label=str(alphas[i]))

plt.xlabel('Number of iterations')
plt.ylabel('Cost J')
plt.legend(loc='upper right')
plt.show()

# Find optimal theta vector with best learning rate
# Assume we already know that alpha = 0.3 is the best learning rate
model = LinearRegression(alpha=0.3, check_conv=True)
theta, _ = model.fit(X_scaled, y)
# print(theta)

```

```

test_data = np.array([1650, 3])
test_data_scaled = (test_data - mu) / sigma
# print(test_data_scaled)
test_data_scaled = np.append([1], test_data_scaled)
# print(test_data_scaled)
print('Predicted price using gradient descent: ',
model.predict(test_data_scaled))

# test_data_scaled.shape = (3, 1)
# print(theta.T.dot(test_data_scaled))

X = np.column_stack((intercept, X))
test_data = np.append([1], test_data)
theta_normal_equation = normal_equation(X, y)
print('Predicted price using normal equation: ',
theta_normal_equation.T.dot(test_data))

def normalize(X):
    # Preprocess data to give std of 1 and mean of 0
    sigma = np.std(X, axis=0) # std
    mu = np.mean(X, axis=0) # mean
    X = (X-mu) / sigma # adjustment
    return X

def normal_equation(X, y):
    theta = inv(X.T.dot(X)).dot(X.T).dot(y)
    return theta

# ----- END PART B ----- #

part_a()
part_b()

```

Part C

```

import numpy as np
import matplotlib.pyplot as plt
from math import log
from numpy.linalg import inv

def newton_method(X, y, theta, tolerance=1e-5, max_iters=15):
    MAX_ITERS = 1
    epoch = 1
    for _ in range(max_iters):

```

```

        H = hessian(X, theta)
        # print(hes)
        g = gradient(X, y, theta)
        # print('Gradient shape: ', g.shape)
        temp = theta - np.dot(inv(H), g)
        if np.sum(abs(theta - temp)) < tolerance:
            print('Convergered at epoch %d' % epoch)
            break
        theta = temp
        epoch += 1
    return theta

def gradient(X, y, theta):
    m = X.shape[0]
    h = hypothesis(X, theta)
    g = (1.0/m) * X.T.dot(h-y)
    return g

def hessian(X, theta):
    m = X.shape[0]
    h = hypothesis(X, theta)
    h.shape = (len(h),)
    H = (1.0/m) * np.dot(np.dot(X.T, np.diag(h)), np.dot(np.diag(1-h), X))
    return H

def sigmoid(z):
    result = 1.0/(1.0+np.exp(-z))
    return result

def cost_function(X, y, theta):
    m = X.shape[0]
    h = hypothesis(X, theta)
    J = (1.0/m) * (-y.dot(np.log(h)) - (1-y).dot(np.log(1-h)))
    return J

def hypothesis(X, theta):
    # print('Hypothesis: ', h.shape)
    return sigmoid(X.dot(theta))

def main():
    X = np.loadtxt('data/cx.dat')
    y = np.loadtxt('data/cy.dat')

    # Get positive and negative indices
    pos = np.nonzero(y)

```

```

neg = np.where(y==0)

# Plot
plt.scatter(X[pos,0], X[pos,1], color='b', marker='+')
plt.scatter(X[neg,0], X[neg,1], color='r', marker='o')
# plt.show()

m = X.shape[0]
X = np.column_stack((np.ones((m, 1)), X))
y.shape = (y.shape[0], 1)
theta = np.zeros((X.shape[1], 1))

theta = newton_method(X, y, theta)
print(theta)

# predict if the student got 20 in exam 1 and 80 in exam 2
test_data = np.array([1, 20, 80])
p = 1 - hypothesis(test_data, theta)
print('Probability that student is not admitted with a score of 20 on
Exam 1 and a score of 80 on Exam 2 is %f' % p)

# Plot decision boundary
min_X = np.min(X[:,1:3])
max_X = np.max(X[:,1:3])
plot_x = [min_X-2, max_X+2]
plot_y = (-1/theta[2])*(theta[1]*plot_x) + theta[0]
plt.plot(plot_x, plot_y)
plt.show()

main()

```

Appendix

Part A: